Architectural Foundations for Ontological Supply Chain Virtual Twins: A Comprehensive Strategic Framework

The contemporary industrial landscape is undergoing a profound shift from traditional data management to semantic intelligence. Central to this transformation is the development of the "Virtual Twin," a digital representation of physical systems that transcends simple data replication by incorporating the underlying logic, relationships, and dynamic behaviors—often referred to as kinetics—of the assets it mirrors.[1] To construct a resilient virtual twin on top of disparate Enterprise Resource Planning (ERP) data, an organization must first establish an ontology. This ontology serves as the conceptual "coordinates," providing a machine-readable blueprint of the organization's world.[3] Unlike a technical database schema, an ontology captures the formal, shared vocabulary of a domain, defining what entities exist, how they are characterized, and the governing rules of their interactions.[4]

The Conceptual Architecture of an Ontology

Building an ontology begins with identifying its core building blocks in the most generic sense, moving away from technical implementation details like graph databases and toward semantic first principles. These blocks facilitate interoperability across systems and allow artificial intelligence to reason over data with the same contextual awareness as a human expert.[6]

Classes and Conceptual Hierarchies

Classes, also referred to as concepts or types, represent the categories of objects within a domain.[3] In a supply chain, a class identifies "what" an object is, such as a Supplier, a Purchase Order, or a Manufacturing Plant.[4] These classes are typically organized into hierarchies through a process known as subsumption, where specific subclasses (e.g., Raw Material) inherit the properties and relationships of more general parent classes (e.g., Product).[6] This taxonomic structure is vital for data integrity, as it ensures that any rule applied to a general category is automatically respected by its specialized members.[10]

Individuals and Ground-Level Instances

Individuals, or instances, represent the actual real-world objects that populate the classes.[8] If "Warehouse" is a class, then "Frankfurt Distribution Center" is an individual.[8] The virtual twin relies on these instances to maintain a live record of physical assets. While the ontology provides the abstract structure, the knowledge graph maps these structures to specific data points from ERP systems, turning the blueprint into a populated "house".[6]

Attributes and Semantic Properties

Attributes, often called data properties, define the characteristics or parameters of a class or individual.[3] These are the qualitative and quantitative values that provide detail, such as the "Lead Time" of a supplier, the "Weight" of a part, or the "Capacity" of a storage location.[4] In an ontological sense, attributes are more than just table columns; they carry metadata regarding their validity, units of measure, and constraints.[12]

Relationships and Object Properties

Relationships, or link types, define how different classes and individuals connect to one another.[3] In a supply chain, these links represent the flow of goods and information: a Supplier *produces* a Material, which is *stored in* a Warehouse, and *shipped via* a Carrier.[4] These relationships are often directional and can possess characteristics such as being inverse (e.g., "Contains" is the inverse of "PartOf") or transitive (e.g., if Part A is PartOf Assembly B, and Assembly B is PartOf Product C, then Part A is PartOf Product C).[8]

Axioms and Logical Constraints

Axioms are the formal assertions and rules that comprise the overall theory of the ontology.[8] They constrain the usage of concepts and relationships to provide real semantics.[10] For example, an axiom might state that "a Shipment cannot be delivered before it has been dispatched" or that "a Production Order must always be associated with exactly one Plant".[10] Axioms allow the virtual twin to perform automated reasoning, identifying contradictions in ERP data that would be invisible in a standard relational database.[5]

Ontology Building Block	**Semantic Function**	**Supply Chain Application Example**
Class	Defines categories or types.	Supplier, Material, LogisticsNode.
Individual	Represents specific instances.	Acme Corp, SKU-402, Long Beach Port.
Attribute	Defines characteristics.	LeadTimeDays, UnitPrice, CurrentStock.
Relationship	Defines connections.	*Supplies*, *Transports*, *AllocatedTo*.
Axiom	Sets logical boundaries.	"StockLevel cannot be less than zero."
Interface	Ensures polymorphism.	Shared behavior for "TransportationAsset."

Palantir’s Ontological Framework: Objects, Links, and Actions

Palantir Foundry approaches ontology building through a distinct methodology that separates the world into semantic elements (which describe the state) and kinetic elements (which describe change).[1] This duality is essential for a virtual twin, as it allows the system to mirror both the static structure and the dynamic operations of the organization.[1]

Semantic Elements: Object and Link Types

In the Palantir model, the foundational layer is composed of Object Types and Link Types. An Object Type is a schema definition of a real-world entity, such as an Aircraft or a Work Order.[3] Each object corresponds to a record in an underlying dataset (e.g., an ERP table), but it is enriched with metadata, value formatting, and security policies.[1] Link Types define the relationships between these objects, functioning as first-class citizens that allow users to navigate the "Search Around" graphs to find dependencies, such as which suppliers are impacted by a specific factory shutdown.[1]

Kinetic Elements: Action Types and Functions

To model the kinetics of an organization, Palantir introduces Action Types and Functions.[1] Action Types define a set of changes that can be applied to objects, properties, and links.[17] For instance, an "Update Delivery Date" action might modify the "ExpectedArrival" property of a Shipment object and simultaneously update the linked "ProductionSchedule".[17] Functions are pieces of code-based logic that allow for the authoring of complex business rules, such as calculating real-time risk scores or optimizing inventory levels based on current market signals.[1]

Advanced Modeling: Interfaces and Shared Properties

Palantir facilitates scalability through Interfaces and Shared Properties. Interfaces allow for polymorphism, enabling consistent modeling of different object types that share a common shape.[1] For example, "Vessel" and "Truck" might both implement a "Carrier" interface, allowing a logistics application to interact with both using the same set of properties like "CurrentLocation".[3] Shared Properties allow for centralized management of metadata across multiple object types, ensuring that a term like "VendorID" is defined and governed consistently throughout the entire ontology.[3]

Industry Standards and Comparative Methodologies

While proprietary platforms like Palantir offer integrated environments, other organizations leverage global standards to build supply chain ontologies. The most prominent of these is the Supply Chain Operations Reference (SCOR) model and its digital evolution, the SCOR Digital Standard (SCOR DS).[10]

The SCOR Digital Standard (SCOR DS) Information Model

SCOR DS provides a unified information model architecture that captures domain knowledge in a formal, machine-readable format.[10] It is organized around primary management processes: Plan, Source, Make, Deliver, Return, and Enable.[19] The SCOR DS information model (SDSIM) uses

a Linked Data technology stack (XML, RDF, OWL) to represent these processes as a hierarchy of classes and relations.[10]
A critical component of SCOR DS is its performance section, which links metrics to specific processes through performance attributes such as Reliability, Responsiveness, Agility, Cost, and Asset Management Efficiency.[10] By modeling these as ontological elements, an organization can diagnose performance gaps by decomposing high-level metrics (e.g., Order Fulfillment Cycle Time) into their constituent operational parts.[19]

The Industrial Ontology Foundry (IOF) and SAREF

The Industrial Ontology Foundry (IOF) is an international collaborative effort to create a coherent set of modular ontologies for the manufacturing and supply chain sectors.[22] These reference ontologies provide a top-level structure that helps organizations avoid logical inconsistencies and gaps in their domain modeling.[23] Similarly, the SAREF (Smart Applications REFerence) ontology emphasizes modularity and extensibility, allowing different stakeholders to specialize concepts like "Device" or "Function" for specific domains while maintaining a common semantic core.[15]

Feature	SCOR DS Approach	Palantir Foundry Approach	SAREF/Academic Approach
Primary Focus	Process-centric benchmarking.	Workflow-centric decision making.	Interoperability & modularity.
Logic Layer	Hierarchical metrics & attributes.	Actions, Functions, & Side Effects.	State machines & state transitions.
Data Format	RDF/OWL (Standardized).	Proprietary Object Storage (Virtual).	RDF/OWL/JSON-LD.
Kinetics	Strategic & Tactical planning.	Operational Action Types.	Formal behavior equations.
Extensibility	Industry-neutral levels 1-3.	Interface-based polymorphism.	Modular extensions (e.g., SAREF4BLDG).

Strategy for Building an Ontology on Top of ERP Data

The strategy for creating a virtual twin begins with the transformation of technical ERP tables into a semantic layer. This process requires a deep understanding of the "semantic gap" between how data is stored in systems like SAP and how it is understood by business users.[25]

Identifying the Domains of Interest

An effective supply chain ontology is typically divided into distinct sectors, each representing a part of the value chain.[9]

- **Upstream (Suppliers):** Dealing with procurement, vendor master data, and raw material sourcing.[9]
- **Internal (Production):** Focused on internal supply chain, operations, and production

orders.[9]

- **Downstream (Clients):** Addressing demand, logistics, and customer fulfillment.[9]
- **Resource Planning (Materials):** Managing materials, stock levels, and item hierarchies.[9]

Mapping ERP Tables to Ontological Objects

Mapping involves translating the cryptic table and field names of an ERP into business-meaningful concepts.[4] For example, in an SAP S/4HANA or ECC system, the EKKO table represents Purchase Order headers, while EKPO contains line items.[26] The ontological strategy would declare a "PurchaseOrder" class and map it to these tables, using the EBELN (Purchase Order Number) field as the unique identifier and linking it to the "Vendor" class via the LFA1 table.[27]

ERP Source Table	Business Meaning	Ontological Class	Key Attributes
LFA1	Vendor Master Data	**Supplier**	Name, Region, LeadTime.
MARA	Material Master Data	**Product**	SKU, MaterialGroup, UnitWeight.
EKKO	PO Header Data	**PurchaseOrder**	OrderDate, TotalValue, VendorID.
EKPO	PO Item Data	**OrderItem**	Quantity, Price, Plant, MaterialID.
MSEG	Material Documents	**Event / Transaction**	Date, MovementType, Quantity.
T001W	Plant Locations	**ManufacturingSite**	Location, Capacity, PlantCode.

The "Ontology-First" Implementation Path

Transitioning to a virtual twin is best accomplished through an "ontology-first" AI strategy.[30] This means starting with the modeling of business concepts and relationships rather than wiring applications directly to data schemas.[30] The roadmap includes:

1. **Defining a Minimal Ontology:** Identify 3–5 core entities (e.g., Supplier, Part, Shipment) and their relationships to create a "Minimum Viable Ontology" (MVO).[7]
2. **Creating Semantic Mappings:** Develop mapping files (YAML, JSON) that tie these concepts to implementation tables in the data warehouse or ERP.[13]
3. **Deploying Ontology-Aware Tools:** Build services that AI agents can use to fetch data, such as get_supplier_performance or list_delayed_shipments, which reason over the ontology rather than raw SQL.[30]

Modeling Kinetics and Behavior on Top of Data

The "kinetics" of a supply chain represent its dynamic life—how it moves, transforms, and reacts to disruptions.[1] In a virtual twin, modeling kinetics involves capturing state transitions and behavioral rules that govern how entities evolve over time.[2]

State Machines and Event-Oriented Modeling

To simulate behavior, the ontology must incorporate a time set and an event set.[31]

- **State Space (S):** The current quantification of an individual's indicators (e.g., a "Shipment" state might be "Dispatched," "InTransit," or "Delivered").[31]
- **Event Set (E):** Instantaneous occurrences that trigger a state change, such as a "Goods Receipt" event firing when a truck arrives at a warehouse.[31]
- **Time Set (T):** The passage of time, which can be discrete or continuous, allowing for the simulation of event lifetimes and delays.[31]

Behavior Equations and Mechanistic Models

Kinetics can also be modeled through mechanistic behavior equations that define the physical principles of the system.[32] For a supply chain, this might involve modeling the "mass balance" of inventory or the "flow rates" of transportation routes.[33] These mechanistic models are often "glass-box" (transparent), allowing for reliable predictions even in scenarios where historical data is sparse.[32]

A general kinetic equation for an inventory state (I) at time t can be expressed as:

$$I(t) = I(t_0) + \int_{t_0}^{t} (\text{Inflow}(r) - \text{Outflow}(r))dr$$

Where:

- **Inflow** is determined by events such as "Purchase Order Received."
- **Outflow** is determined by events such as "Sales Order Shipped."

By embedding these equations into the ontology, the virtual twin can simulate future states and identify potential bottlenecks or stock-outs before they occur.[2]

Discrete-Event Simulation (DES) for Virtual Twins

Discrete-Event Simulation (DES) is a cornerstone of kinetic modeling.[35] It mimics the operation of a system as a discrete sequence of events, which is ideal for supply chains characterized by variability and resource constraints.[35] DES provides the foundational technology to not only analyze the system but to construct the "digital replica" that serves as the heart of the virtual twin.[35] When a disruption occurs—such as a port closure—the DES engine uses the ontology's relationships to simulate the impact across the network and test alternative "what-if" scenarios.[35]

Using AI to Mine Insights and Drive Intelligence

Once the ontology and kinetics are in place, artificial intelligence can be leveraged to mine insights and automate complex decision-making processes. The ontology acts as the "semantic foundation" that allows AI to reason with the contextual awareness of a human operator.[6]

Reasoning via Knowledge Graph Traversal

In an ontological virtual twin, AI agents do not just look at dashboards; they traverse the relationships defined in the knowledge graph to find root causes.[4] If a customer order is at risk, an agent follows the edges backward from the "Order" to the "Shipment," then to the "Vehicle," and finally to a "Risk Event" like a traffic delay or weather disruption.[4] This "graph traversal" allows for instantaneous root cause analysis that would take humans hours to manualy trace through ERP silos.[4]

Ontology Augmented Generation (OAG)

Large Language Models (LLMs) can be integrated with the ontology through Ontology Augmented Generation (OAG).[39] In this paradigm, the ontology guides the LLM in an iterative, chain-of-thought process to extract semantic information from unstructured open data—such as news reports or government filings.[39] For example, an LLM can be prompted to identify "Natural Hazard" events in news articles and map them directly into the knowledge graph as "Risk" individuals linked to specific "Suppliers".[39]

Predictive and Prescriptive Insights

By analyzing the "traces" of events stored in the ontology, AI models can move from descriptive analytics (what happened) to predictive and prescriptive analytics (what will happen and what should we do).[30]

- **Anomaly Detection:** AI algorithms run against the graph structure to detect patterns that deviate from the "normative" axioms of the ontology.[4]
- **Scenario Optimization:** The virtual twin can temporarily modify "edges" in the graph (e.g., switching to an alternate carrier) and simulate the outcomes to recommend the optimal path forward.[4]
- **Decision Synchronization:** AI helps align choices across the entire supply chain network, ensuring that when a procurement officer changes a supplier, the impact is immediately understood and compensated for in the logistics and manufacturing layers.[41]

Operationalizing the Virtual Twin: Integration and Scale

Building the virtual twin requires more than just an ontology; it requires a platform capable of real-time data ingestion, transformation, and service delivery.[24]

Architecture for Real-Time Synchronization

A resilient virtual twin architecture is often based on microservices and containerization.[24] Data from the physical system is ingested through protocols like MQTT or OPC-UA and routed to a "broker" that manages the real-time entities in the knowledge graph.[24] This ensures that simulation engines and analytics tools always have access to fresh, well-organized data.[24]

Secure and Interoperable Data Spaces

For organizations operating in complex multi-partner networks, the use of "Data Space" technology is emerging as a critical standard.[40] This approach incorporates semantic aspect meta-models to ensure that data exchanged between partners retains its meaning and adheres to specific policies (permissions, duties, and prohibitions).[40] This is vital for maintaining the "Open World Assumption," where new partners and data points can be integrated into the twin without breaking the existing ontological structure.[10]

Maintaining and Evolving the Ontology

The ontology is not a static document but a living model that must evolve alongside the physical supply chain.[1] Maintenance involves identifying and correcting defects, accommodating new business requirements (e.g., adding "Sustainability" metrics for ESG compliance), and managing the versioning of action types and functions to ensure consistency over time.[7]

Implementation Phase	**Strategic Objective**	**Key Deliverables**
Phase 1: Conceptualization	Define the shared vocabulary.	MVO (Minimum Viable Ontology), Sector Definitions.
Phase 2: Semantic Mapping	Bridge the gap to ERP data.	Mapping Layers (SQL to Ontology), Data Cleaning Rules.
Phase 3: Kinetic Modeling	Introduce movement & state.	State Machines, Discrete-Event Simulation Models.
Phase 4: Intelligence Layer	Enable AI reasoning & insights.	OAG Pipelines, Knowledge Graph AI Agents.
Phase 5: Decision Support	Power end-user workflows.	Object Views, Action-Log Auditing, "What-If" Analysis.

Conclusion: The Strategic Value of an Ontological Virtual Twin

The development of an ontological virtual twin represents the pinnacle of modern supply chain engineering. By moving beyond technical data storage and into the realm of semantic modeling, organizations can create a digital replica that truly "understands" the physical world it represents.[1] The generic building blocks of classes, relations, attributes, and axioms provide the necessary coordinates to translate messy ERP data into a structured knowledge base.[4] Palantir's methodology of combining semantic "Objects" with kinetic "Actions" offers a powerful template for building such a system, while international standards like SCOR DS ensure that the resulting model is benchmarkable and interoperable.[1] When combined with artificial intelligence and discrete-event simulation, this ontological foundation allows for the modeling of complex kinetics, enabling a level of predictive and prescriptive insight that was previously unattainable.[4]

Ultimately, the goal of investing in an ontology for a supply chain virtual twin is to facilitate better, faster decision-making.[1] In an era of constant disruption, from geopolitical conflicts to climate-driven disasters, the ability to "stress-test" the organization in a digital environment before committing real-world resources is no longer a luxury—it is a strategic necessity for resilience and competitive advantage.[38]

Works cited

1. Overview • Ontology • Palantir, accessed February 28, 2026, https://palantir.com/docs/foundry/ontology/overview/
2. From Mathematical Modeling and Simulation to Digital Twins: Bridging Theory and Digital Realities in Industry and Emerging Technologies - MDPI, accessed February 28, 2026, https://www.mdpi.com/2076-3417/15/16/9213
3. Core concepts - Palantir, accessed February 28, 2026, https://palantir.com/docs/foundry/ontology/core-concepts/
4. How Graph Databases & Ontology Will Rewire Supply Chains Forever - Medium, accessed February 28, 2026, https://medium.com/@emails.chandra04/how-graph-databases-ontology-will-rewire-supply-chains-forever-917d0a26b281
5. Why You Should Consider Ontology Modeling for AI-Driven Digital Twins | by Elisha Miller, accessed February 28, 2026, https://medium.com/timbr-ai/why-you-should-consider-ontology-modeling-for-ai-driven-digital-twins-c36a2319e22c
6. Building Ontology-Driven Intelligence for Industrial AI Agents - HiveMQ, accessed February 28, 2026, https://www.hivemq.com/resources/building-ontology-driven-intelligence-for-industrial-ai-agents
7. Supply Chain's AI Building Blocks: Knowledge Graphs vs Ontologies - Zero100, accessed February 28, 2026, https://zero100.com/supply-chains-ai-building-blocks-knowledge-graphs-vs-ontologies/
8. Ontology components - Wikipedia, accessed February 28, 2026, https://en.wikipedia.org/wiki/Ontology_components
9. Ontology-based Knowledge Graph for Supply ... - IPMU2024 Lisboa, accessed February 28, 2026, https://ipmu2024.inesc-id.pt/files/paper_2195.pdf
10. The ASCM SCOR Digital Standard Information Model, accessed February 28, 2026, https://www.ascm.org/globalassets/ascm_website_assets/docs/scor/information-model-scor-digital-standard-2025.pdf
11. Object and link types • Types reference - Palantir, accessed February 28, 2026,

https://palantir.com/docs/foundry/object-link-types/type-reference/
12. Object and link types • Properties • Overview - Palantir, accessed February 28, 2026, https://palantir.com/docs/foundry/object-link-types/properties-overview/
13. Semantic Layers in 2025: A Data Leaders Playbook - Coalesce, accessed February 28, 2026, https://coalesce.io/data-insights/semantic-layers-2025-catalog-owner-data-leader-playbook/
14. Ontology Object Properties are Data Model Associative Entities - not Relationships, accessed February 28, 2026, https://fib-dm.com/ontology-object-property-data-model-associative-entities/
15. SAREF Design Principles - SAREF ontology - ETSI, accessed February 28, 2026, https://saref.etsi.org/principles.html
16. Ontologies in Digital Twins: A Systematic Literature Review - UvA-DARE (Digital Academic Repository), accessed February 28, 2026, https://pure.uva.nl/ws/files/154600806/Ontologies_in_Digital_Twins.pdf
17. Action types • Overview - Palantir, accessed February 28, 2026, https://palantir.com/docs/foundry/action-types/overview/
18. Action types • Side effects • Overview - Palantir, accessed February 28, 2026, https://palantir.com/docs/foundry/action-types/side-effects-overview/
19. SCOR Digital Standard - ASCM, accessed February 28, 2026, https://www.ascm.org/globalassets/ascm_website_assets/docs/intro-and-front-matter-scor-digital-standard.pdf
20. SUPPLY CHAIN DYNAMICS IN THE SCOR MODEL - A SIMULATION MODELING APPROACH, accessed February 28, 2026, https://informs-sim.org/wsc12papers/includes/files/con245.pdf
21. (PDF) Ontology-Based Semantic Model of Supply Chains for Modeling and Simulation in Distributed Environment - ResearchGate, accessed February 28, 2026, https://www.researchgate.net/publication/312559687_Ontology-Based_Semantic_Model_of_Supply_Chains_for_Modeling_and_Simulation_in_Distributed_Environment
22. Towards Ontologizing a Digital Twin Framework for Manufacturing - NIST, accessed February 28, 2026, https://www.nist.gov/publications/towards-ontologizing-digital-twin-framework-manufacturing
23. Modeling a Supply Chain Reference Ontology Based on a Top-Level Ontology | NIST, accessed February 28, 2026, https://www.nist.gov/publications/modeling-supply-chain-reference-ontology-based-top-level-ontology
24. Full article: Development and demonstration of a digital twin platform leveraging ontologies and data-driven simulation models - Taylor & Francis, accessed February 28, 2026, https://www.tandfonline.com/doi/full/10.1080/19401493.2025.2504005
25. Approach to database-to-ontology mapping . | Download Scientific Diagram - ResearchGate, accessed February 28, 2026,

https://www.researchgate.net/figure/Approach-to-database-to-ontology-mapping_fig3_51025142
26. My First Steps into Mastering SAP's Data Models | Towards Data Science, accessed February 28, 2026, https://towardsdatascience.com/my-first-steps-into-mastering-saps-data-models-4d20ad2485f2/
27. SAP: EKKO and EKPO tables - EIP International Business School, accessed February 28, 2026, https://eiposgrados.com/eng/sap-consultant/sap-ekko-and-ekpo-tables/
28. SAP Table EKPO - Purchasing Document Item - LeanX, accessed February 28, 2026, https://leanx.eu/sap/table/ekpo/
29. EKKO & EKPO Table in SAP - Purchase Order (PO) Tables in SAP - SAP Tutorials, accessed February 28, 2026, https://www.saponlinetutorials.com/ekko-ekpo-table-in-sap/
30. Ontology-First AI Architecture: Definition, Components, Examples ..., accessed February 28, 2026, https://atlan.com/know/ontology-first-ai-architecture/
31. DeMO: An Ontology for Discrete-event Modeling and Simulation - PMC, accessed February 28, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC3423901/
32. Automated Generation of Mechanistic Models for Chemical Process Digital Twins using Reinforcement Learning - Part II - ChemRxiv, accessed February 28, 2026, https://chemrxiv.org/doi/pdf/10.26434/chemrxiv-2025-83x6v
33. Ontology-Based Process Modelling-with Examples of Physical ..., accessed February 28, 2026, https://www.mdpi.com/2227-9717/9/4/592
34. Model management to support systems engineering workflows using ontology-based knowledge graphs - arXiv.org, accessed February 28, 2026, https://arxiv.org/html/2512.09596v1
35. Digital Twins and Discrete-Event Simulation: A Transformational Partnership - Simio, accessed February 28, 2026, https://www.simio.com/digital-twins-and-discrete-event-simulation-a-transformational-partnership/
36. Enhancing internal supply chain management in manufacturing through a simulation-based digital twin platform - IRIS, accessed February 28, 2026, https://iris.unisalento.it/bitstream/11587/539706/1/17_Journal.pdf
37. Supply Chain Digital Twins - anyLogistix, accessed February 28, 2026, https://www.anylogistix.com/features/supply-chain-digital-twins/
38. The Complete Guide to Supply Chain Digital Twins - Coupa, accessed February 28, 2026, https://www.coupa.com/blog/the-complete-guide-to-supply-chain-digital-twins/
39. Empowering Supply Chain Risk Monitoring with ... - CEUR-WS.org, accessed February 28, 2026, https://ceur-ws.org/Vol-4085/paper4.pdf
40. Supply Chain Data Analytics for Digital Twins: A Comprehensive ..., accessed February 28, 2026, https://www.mdpi.com/2076-3417/15/12/6939
41. The $31 Billion Question: What AI and Machine Learning Really Do for Supply Chains | Walton College | University of Arkansas, accessed February 28, 2026, https://walton.uark.edu/insights/posts/the-31-billion-dollar-question-what-ai-and

-machine-learning-really-do-for-supply-chains.php
42. Digital Twins in Supply Chain Management, accessed February 28, 2026, https://www.digitaltwinconsortium.org/2024/06/digital-twins-in-supply-chain-management/